

```
In [49]: 1 import numpy as np
          2 import pandas as pd
          3 import seaborn as sns
          4 import matplotlib.pyplot as plt
          5 from sklearn.model_selection import train_test_split
          6 from sklearn.linear_model import LinearRegression
          7 from sklearn import metrics
          8 from sklearn import preprocessing,svm
          9
```

```
In [50]: 1 df=pd.read_csv(r"C:\Users\P. VIJAY KUMAR\Downloads\Advertising.csv")
          2 df
```

Out[50]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

```
In [51]: 1 df.head()
```

Out[51]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [52]: 1 df.tail()
```

Out[52]:

	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

In [53]: 1 df.describe()

Out[53]:

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

In [54]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio       200 non-null    float64
2    Newspaper   200 non-null    float64
3    Sales       200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

In [55]: 1 df.shape

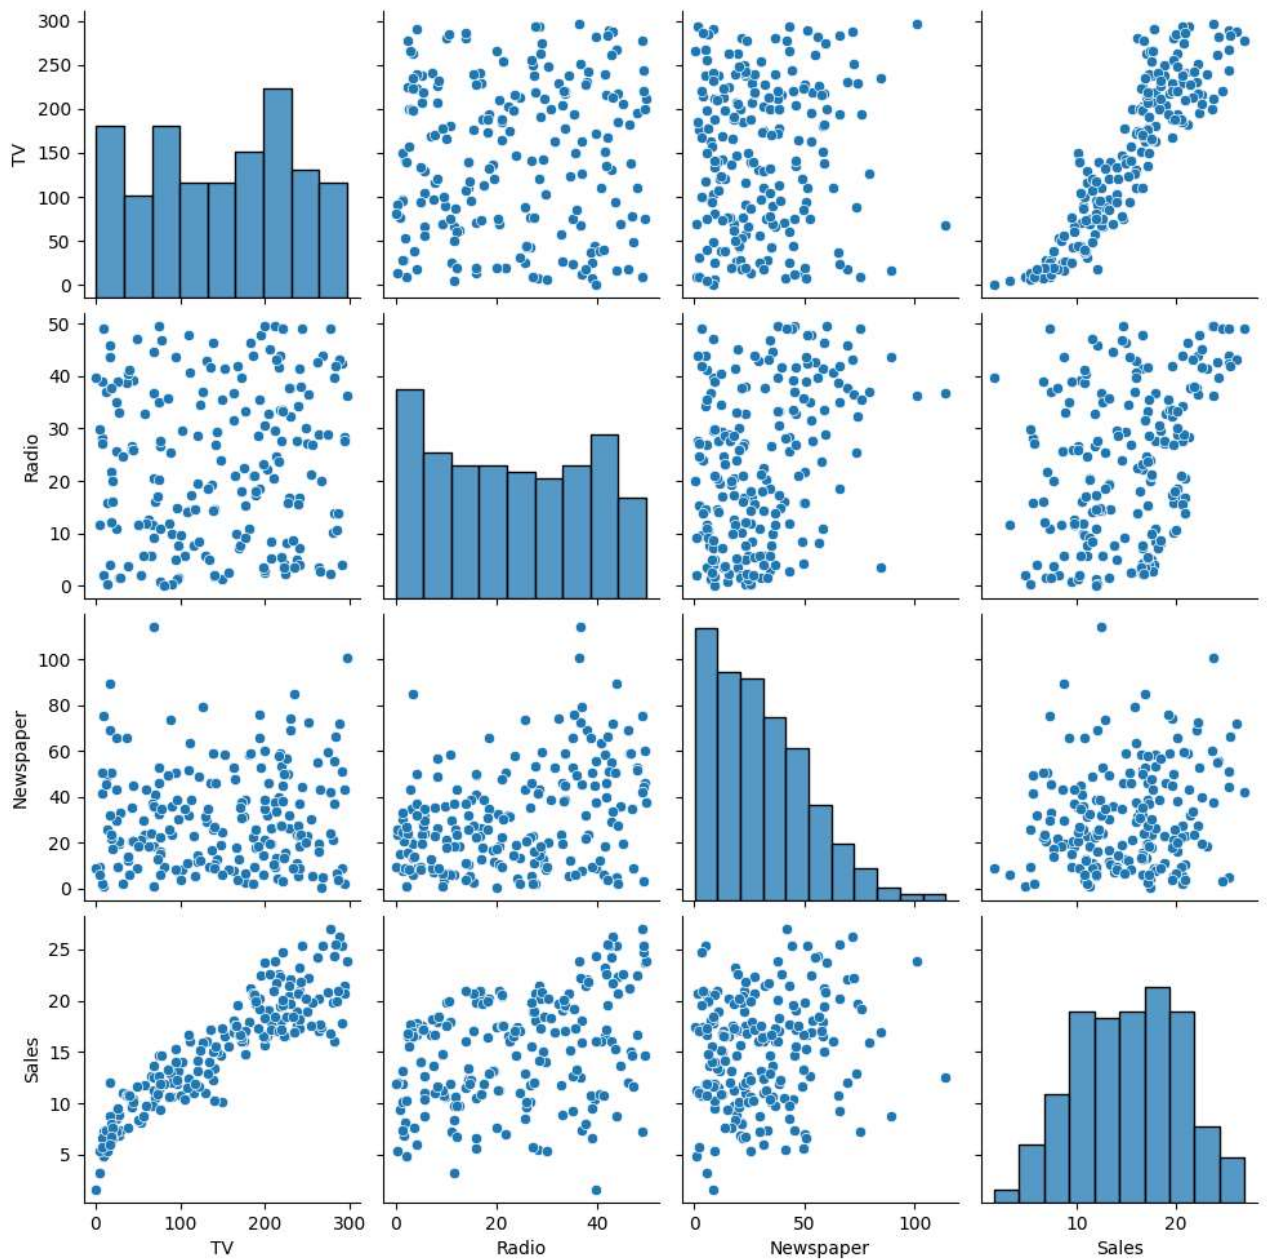
Out[55]: (200, 4)

In [56]: 1 df.columns

Out[56]: Index(['TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')


```
In [57]: 1 sns.pairplot(df)
```

```
Out[57]: <seaborn.axisgrid.PairGrid at 0x212af247910>
```




```
In [58]: 1 sns.displot(df['Sales'])
```

```
Out[58]: <seaborn.axisgrid.FacetGrid at 0x212af6dfb20>
```



```
In [59]: 1 addf=df[['TV', 'Radio', 'Newspaper', 'Sales']]
2 sns.heatmap(addf.corr())
3
```

```
Out[59]: <Axes: >
```



```
In [60]: 1 X=addf[['TV', 'Radio', 'Newspaper']]
        2 y=df['Sales']
```

```
In [61]: 1 from sklearn.model_selection import train_test_split
        2 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=101)
        3 from sklearn.linear_model import LinearRegression
        4 lm=LinearRegression()
        5 lm.fit(X_train,y_train)
        6 print(lm.intercept_)
```

4.681232151484295

```
In [62]: 1 coeff_df=pd.DataFrame(lm.coef_,X.columns,columns=['coefficient'])
        2 coeff_df
        3
```


Out[62]:

	coefficient
TV	0.054930
Radio	0.109558
Newspaper	-0.006194

```
In [63]: 1 predictions=lm.predict(X_test)
```


```
In [64]: 1 plt.scatter(y_test,predictions)
```

Out[64]: <matplotlib.collections.PathCollection at 0x212ae8a2b60>



```
In [65]: 1 sns.displot((y_test,predictions),bins=50)#without semicolon
```

```
Out[65]: <seaborn.axisgrid.FacetGrid at 0x212af639ed0>
```



```
In [66]: 1 from sklearn import metrics
2 print('MAE:',metrics.mean_absolute_error(y_test,predictions))
3 print('MSE:',metrics.mean_squared_error(y_test,predictions))
4 print('MAE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```


MAE: 1.3731200698367851

MSE: 2.8685706338964967

MAE: 1.6936855180040056


```
In [67]: 1 sns.lmplot(x="Newspaper",y="Sales",data=df,order=2)
```

```
Out[67]: <seaborn.axisgrid.FacetGrid at 0x212b01c2260>
```




```
In [68]: 1 sns.lmplot(x="TV",y="Sales",data=df,order=2)
```

```
Out[68]: <seaborn.axisgrid.FacetGrid at 0x212b02a87c0>
```



```
In [69]: 1 sns.lmplot(x="Radio",y="Sales",data=df,order=2)
        2
```

Out[69]: <seaborn.axisgrid.FacetGrid at 0x212b039e050>



```
In [70]: 1 df.fillna(method='ffill',inplace=True)
```

```
In [71]: 1 df.fillna(method='ffill',inplace=True)
```

```
In [72]: 1 regr=LinearRegression()
        2
```

```
In [73]: 1 x=np.array(df['TV']).reshape(-1,1)
        2 y=np.array(df['Sales']).reshape(-1,1)
        3 df.dropna(inplace=True)
```


```
In [74]: 1 X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
        2 regr.fit(X_train,y_train)
        3 regr.fit(X_train,y_train)
        4
```

Out[74]:

```
LinearRegression
LinearRegression()
```




```
In [75]: 1 y_pred=regr.predict(X_test)
2 plt.scatter(X_test,y_test,color='y')
3 plt.plot(X_test,y_pred,color='b')
4 plt.show()
```



```
In [76]: 1 sns.pairplot(df,x_vars=['TV', 'Radio', 'Newspaper'],y_vars='Sales',height=7,aspect=0.7,kind='reg')
```

Out[76]: <seaborn.axisgrid.PairGrid at 0x212b04876d0>




```
In [77]: 1 #accuracy
2 regr=LinearRegression()
3 regr.fit(X_train,y_train)
4 regr.fit(X_train,y_train)
5 print(regr.score(X_test,y_test))
6
```

0.797139039213951

```
In [78]: 1 from sklearn.linear_model import Lasso,Ridge
2 from sklearn.preprocessing import StandardScaler
3
```

```
In [79]: 1 ddf=df[['TV', 'Radio', 'Newspaper', 'Sales']]
```

```
In [80]: 1 df.drop(columns = ["Radio", "Newspaper"], inplace = True)
2 sns.pairplot(df)
3 df.Sales=np.log(df.Sales)
4
```



```
In [81]: 1 features=df.columns[0:2]
2 target=df.columns[-1]
3 X=df[features].values
4 y=df[target].values
5 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=17)
6 print("The dimension of X_train is {}".format(X_train.shape))
7 print("The dimension of X_test is {}".format(X_test.shape))
8 scaler=StandardScaler()
9 X_train=scaler.fit_transform(X_train)
10 X_test=scaler.transform(X_test)
```

The dimension of X_train is (140, 2)

The dimension of X_test is (60, 2)

```
In [82]: 1 #Linear regression model
2 regr=LinearRegression()
3 regr.fit(X_train,y_train)
4 actual=y_test #actual value
5 train_score_regr=regr.score(X_train,y_train)
6 test_score_regr=regr.score(X_test,y_test)
7 print("\nLinear model:\n")
8 print("The train score for Linear model is {}".format(train_score_regr))
9 print("The test score for Linear model is {}".format(test_score_regr))
10
```

Linear model:

The train score for Linear model is 1.0
The test score for Linear model is 1.0

```
In [83]: 1 #ridge regression model
2 ridgeReg=Ridge(alpha=10)
3 ridgeReg.fit(X_train,y_train)
4 #train and test score for ridge regression
5 train_score_ridge=ridgeReg.score(X_train,y_train)
6 test_score_ridge=ridgeReg.score(X_test,y_test)
7 print("\nRidge model:\n")
8 print("The train score for ridge model is {}".format(train_score_ridge))
9 print("The test score for ridge model is {}".format(test_score_ridge))
10
```

Ridge model:

The train score for ridge model is 0.9902871391941609
The test score for ridge model is 0.984426628514122


```
In [84]: 1 #using the linear cv model for ridge regression
2 from sklearn.linear_model import RidgeCV
3 #ridge cross validation
4 ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
5 #score
6 print(ridge_cv.score(X_train,y_train))
7 print(ridge_cv.score(X_test,y_test))
```

0.999999999997627
0.9999999999962466

```
In [85]: 1 #using the linear cv model for lasso regression
2 from sklearn.linear_model import LassoCV
3 #lasso cross validation
4 lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y_train)
5 #score
6 print(lasso_cv.score(X_train,y_train))
7 print(lasso_cv.score(X_test,y_test))
8
```

0.9999999343798134
0.9999999152638072

```
In [87]: 1 plt.figure(figsize=(10,10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',
3          label=r'Ridge;\alpha=10$',zorder=7)
4 plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize=7,
5          color='green',label='Linear Regression')
6 plt.xticks(rotation=90)
7 plt.legend()
8 plt.show()
```



```
In [88]: 1 #ridge regression  
2 plt.figure(figsize=(10,10))  
3 sns.heatmap(ddf.corr(),annot=True)
```

Out[88]: <Axes: >




```
In [89]: 1 #Lasso regression model
2 lassoReg=Lasso(alpha=10)
3 lassoReg.fit(X_train,y_train)
4 #train and test score for ridge regression
5 train_score_lasso=lassoReg.score(X_train,y_train)
6 test_score_lasso=lassoReg.score(X_test,y_test)
7 print("\nLasso model:\n")
8 print("The train score for lasso model is {}".format(train_score_lasso))
9 print("The test score for lasso model is {}".format(test_score_lasso))
```

Lasso model:

The train score for lasso model is 0.0


The test score for lasso model is -0.0042092253233847465

```
In [90]: 1 plt.figure(figsize=(10,10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',
3          label=r'Ridge;\alpha=10$',zorder=7)
4 plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize=7,
5          color='green',label='Linear Regression')
6 plt.xticks(rotation=90)
7 plt.legend()
8 plt.show()
```



```
In [91]: 1 pd.Series(lassoReg.coef_, features).sort_values(ascending=True).plot(kind="bar")
```

Out[91]: <Axes: >




In [92]:

```

1 #plot size
2 plt.figure(figsize=(10,10))
3 #add plot for ridge regression
4 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,
5          color='red',label=r'Ridge;\alpha=10$',zorder=7)
6 #add plot for Lasso regression
7 plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,
8          color='green',label=r'Lasso;\alpha=10$',zorder=7)
9 #add plot for linear model
10 plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize=7,
11          color='b',label=r'LinearRegression')
12 #rotate axis
13 plt.xticks(rotation=90)
14 plt.legend()
15 plt.title("Comparison of Ridge,Lasso and Linear regression models")
16 plt.show()

```



In []:

1